

```
In [2]: import pandas as pd
import numpy as np
import seaborn as sns
from matplotlib import pyplot as plt
from sklearn import preprocessing

import warnings
warnings.filterwarnings('ignore')
```

```
In [3]: df = pd.read_csv('scaler_apollo_hospitals.csv')
```

```
In [4]: df.shape
#shape of data
```

```
Out[4]: (1338, 8)
```

```
In [5]: df.head(10)
```

```
Out[5]:
```

	Unnamed: 0	age	sex	smoker	region	viral load	severity level	hospitalization charges
0	0	19	female	yes	southwest	9.30	0	42212
1	1	18	male	no	southeast	11.26	1	4314
2	2	28	male	no	southeast	11.00	3	11124
3	3	33	male	no	northwest	7.57	0	54961
4	4	32	male	no	northwest	9.63	0	9667
5	5	31	female	no	southeast	8.58	0	9392
6	6	46	female	no	southeast	11.15	1	20601
7	7	37	female	no	northwest	9.25	3	18204
8	8	37	male	no	northeast	9.94	2	16016
9	9	60	female	no	northwest	8.61	0	72308

```
In [6]: df.info()
#data type of all attributes
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            1338 non-null   int64
1   age                   1338 non-null   int64
2   sex                   1338 non-null   object
3   smoker                1338 non-null   object
4   region                1338 non-null   object
5   viral load            1338 non-null   float64
6   severity level        1338 non-null   int64
7   hospitalization charges 1338 non-null   int64
dtypes: float64(1), int64(4), object(3)
memory usage: 83.8+ KB
```

```
In [7]: df['region'].value_counts()
#statistical summary of region can be inferred from this
```

```
Out[7]: southeast    364
southwest    325
northwest    325
northeast    324
Name: region, dtype: int64
```

```
In [8]: # conversion of categorical attributes to 'category'

#df = pd.get_dummies(df, columns=["region"])
#one hot encoding for the region would make it more complex but can be done
```

```
#since sex and smoker are binary we can do the below encoding
df['sex'] = df['sex'].map({'female':1, 'male':0})
df['smoker'] = df['smoker'].map({'yes':1, 'no':0})
```

In [9]:

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 8 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Unnamed: 0            1338 non-null   int64
 1   age                   1338 non-null   int64
 2   sex                   1338 non-null   int64
 3   smoker                1338 non-null   int64
 4   region                1338 non-null   object
 5   viral load            1338 non-null   float64
 6   severity level        1338 non-null   int64
 7   hospitalization charges 1338 non-null   int64
dtypes: float64(1), int64(6), object(1)
memory usage: 83.8+ KB
```

In [10]:

```
df.columns
```

Out[10]: Index(['Unnamed: 0', 'age', 'sex', 'smoker', 'region', 'viral load', 'severity level', 'hospitalization charges'], dtype='object')

In [11]:

```
df.rename( columns={'Unnamed: 0':'Index'}, inplace=True )
```

In [12]:

```
#Missing Value Detection
df_nv = df.isna()
```

In [13]:

```
df_nv.value_counts()

#As we can see NO Null values
```

Out[13]:

	Index	age	sex	smoker	region	viral load	severity level	hospitalization charges	
False	False	False	False	False	False	False	False	False	1338

dtype: int64

In [14]:

```
#statistical summary
df.describe()
```

Out[14]:

	Index	age	sex	smoker	viral load	severity level	hospitalization charges
count	1338.000000	1338.000000	1338.000000	1338.000000	1338.000000	1338.000000	1338.000000
mean	668.500000	39.207025	0.494768	0.204783	10.221233	1.094918	33176.058296
std	386.391641	14.049960	0.500160	0.403694	2.032796	1.205493	30275.029296
min	0.000000	18.000000	0.000000	0.000000	5.320000	0.000000	2805.000000
25%	334.250000	27.000000	0.000000	0.000000	8.762500	0.000000	11851.000000
50%	668.500000	39.000000	0.000000	0.000000	10.130000	1.000000	23455.000000
75%	1002.750000	51.000000	1.000000	0.000000	11.567500	2.000000	41599.500000
max	1337.000000	64.000000	1.000000	1.000000	17.710000	5.000000	159426.000000

In [15]:

```
df.nunique()
```

Out[15]:

Index	1338
age	47
sex	2
smoker	2
region	4
viral load	462
severity level	6
hospitalization charges	1320

dtype: int64

```
In [16]: df['severity level'].value_counts()
```


```
Out[16]: 0    574
1    324
2    240
3    157
4     25
5     18
Name: severity level, dtype: int64
```

```
In [17]: #Univariate Analysis

#Categorical variables are : sex, smoker, region, severity load


#Sex
sns.countplot(x="sex", data=df)
```

```
Out[17]: <AxesSubplot:xlabel='sex', ylabel='count'>
```




```
In [18]: sns.countplot(x="sex", hue = "severity level", data=df)
```

```
Out[18]: <AxesSubplot:xlabel='sex', ylabel='count'>
```



```
In [19]: sns.countplot(x="sex", hue = "smoker", data=df)
```

```
Out[19]: <AxesSubplot:xlabel='sex', ylabel='count'>
```






In [20]:

```
#smoker
sns.countplot(x="smoker", data=df)
```

Out[20]:

<AxesSubplot:xlabel='smoker', ylabel='count'>

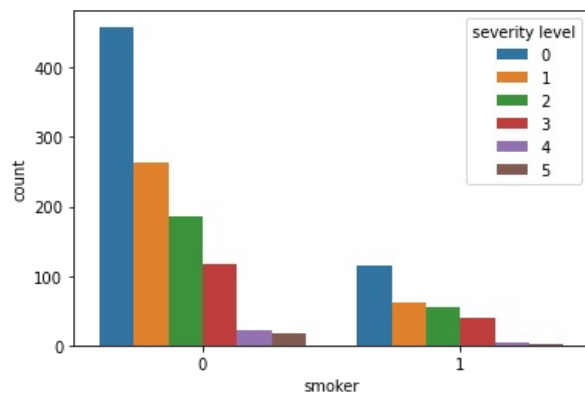


In [21]:

```
sns.countplot(x="smoker", hue = "severity level", data=df)
```

Out[21]:

<AxesSubplot:xlabel='smoker', ylabel='count'>




In [22]:

```
#region
sns.countplot(x="region", data=df)
```

Out[22]:


<AxesSubplot:xlabel='region', ylabel='count'>



In [23]:


```
ssn = df['region'].value_counts()
plt.style.use('seaborn')
plt.figure(figsize = (10, 8))
plt.pie(ssn.values, labels = ssn.index, autopct = '%1.1f%%')
```

```
plt.title('Region Distribution', fontdict = {'fontname' : 'Monospace', 'fontsize' : 30, 'fontweight' : 'bold'})
plt.legend()
plt.axis('equal')
plt.show()
```




```
In [24]: #severity load
sns.countplot(x="severity level", data=df)
```


```
Out[24]: <AxesSubplot:xlabel='severity level', ylabel='count'>
```



```
In [25]: sns.countplot(x="severity level", hue = "smoker", data=df)
```


```
Out[25]: <AxesSubplot:xlabel='severity level', ylabel='count'>
```






```
In [25]: sns.countplot(x="severity level", hue = "sex", data=df)
```

```
Out[25]: <AxesSubplot:xlabel='severity level', ylabel='count'>
```



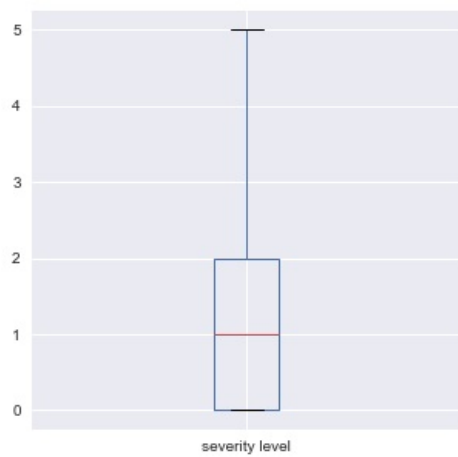
```
In [26]: ssn = df['severity level'].value_counts()
plt.style.use('seaborn')
plt.figure(figsize = (10, 8))
plt.pie(ssn.values, labels = ssn.index, autopct = '%1.1f%')
plt.title('Severity Level Distribution', fontdict = {'fontname' : 'Monospace', 'fontsize' : 30, 'fontweight' : 'bold'})
plt.legend()
plt.axis('equal')
plt.show()
```

Severity Level Distribution



```
In [27]: cat_cols = ['severity level']
df[cat_cols].boxplot(figsize=(5,5))
```

```
Out[27]: <AxesSubplot:>
```




In []:

In [28]: *#Continuous variables are : age, viral load, hospitalization charges*

```
#age
sns.displot( df['age'], kde = True)
```


Out[28]: <seaborn.axisgrid.FacetGrid at 0x2c342952fd0>



In [29]:

```
fig, ax = plt.subplots(figsize=(18, 5))
sns.countplot(x= 'age', data = df, ax= ax)
```


Out[29]: <AxesSubplot:xlabel='age', ylabel='count'>



In [30]:


```
fig, ax = plt.subplots(figsize=(18, 5))
sns.histplot(data = df, x = 'age', ax=ax)
```

Out[30]: <AxesSubplot:xlabel='age', ylabel='Count'>




In [31]: `#viral load`
`sns.displot(df['viral load'], kde= True)`

Out[31]: <seaborn.axisgrid.FacetGrid at 0x2c34295f040>



In [32]: `fig, ax = plt.subplots(figsize=(18, 5))`
`sns.countplot(x= 'viral load', data = df, ax= ax)`


Out[32]: <AxesSubplot:xlabel='viral load', ylabel='count'>



In [33]: `fig, ax = plt.subplots(figsize=(18, 5))`
`sns.histplot(data = df, x = 'viral load', ax=ax)`


Out[33]: <AxesSubplot:xlabel='viral load', ylabel='Count'>






```
In [34]: #hospitalization charges
sns.displot(df['hospitalization charges'], kde=True, aspect = 2)
```

```
Out[34]: <seaborn.axisgrid.FacetGrid at 0x2c342b44a60>
```




```
In [35]: fig, ax = plt.subplots(figsize=(18, 5))
sns.countplot(x= 'hospitalization charges', data = df, ax= ax)
```


```
Out[35]: <AxesSubplot:xlabel='hospitalization charges', ylabel='count'>
```



```
In [36]: fig, ax = plt.subplots(figsize=(18, 5))
sns.histplot(data = df, x = 'hospitalization charges', ax=ax)
```

```
Out[36]: <AxesSubplot:xlabel='hospitalization charges', ylabel='Count'>
```






```
In [37]: #Bivariate analysis

#Heatmap

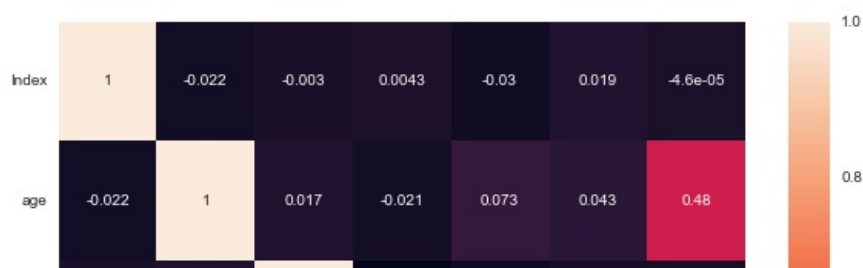
#Since variables aren't normally distributed we do not consider Pearson correlation
fig, ax = plt.subplots(figsize=(10, 10))
Var_Corr = df.corr(method = 'pearson')
sns.heatmap(Var_Corr, xticklabels=Var_Corr.columns, yticklabels=Var_Corr.columns, annot=True, ax=ax)
```


Out[37]: <AxesSubplot:>



```
In [38]: fig, ax = plt.subplots(figsize=(10, 10))
Var_Corr = df.corr(method = 'kendall')
sns.heatmap(Var_Corr, xticklabels=Var_Corr.columns, yticklabels=Var_Corr.columns, annot=True, ax=ax)
```

Out[38]: <AxesSubplot:>






```
In [39]: fig, ax = plt.subplots(figsize=(10, 10))
Var_Corr = df.corr(method = 'spearman')
sns.heatmap(Var_Corr, xticklabels=Var_Corr.columns, yticklabels=Var_Corr.columns, annot=True, ax=ax)

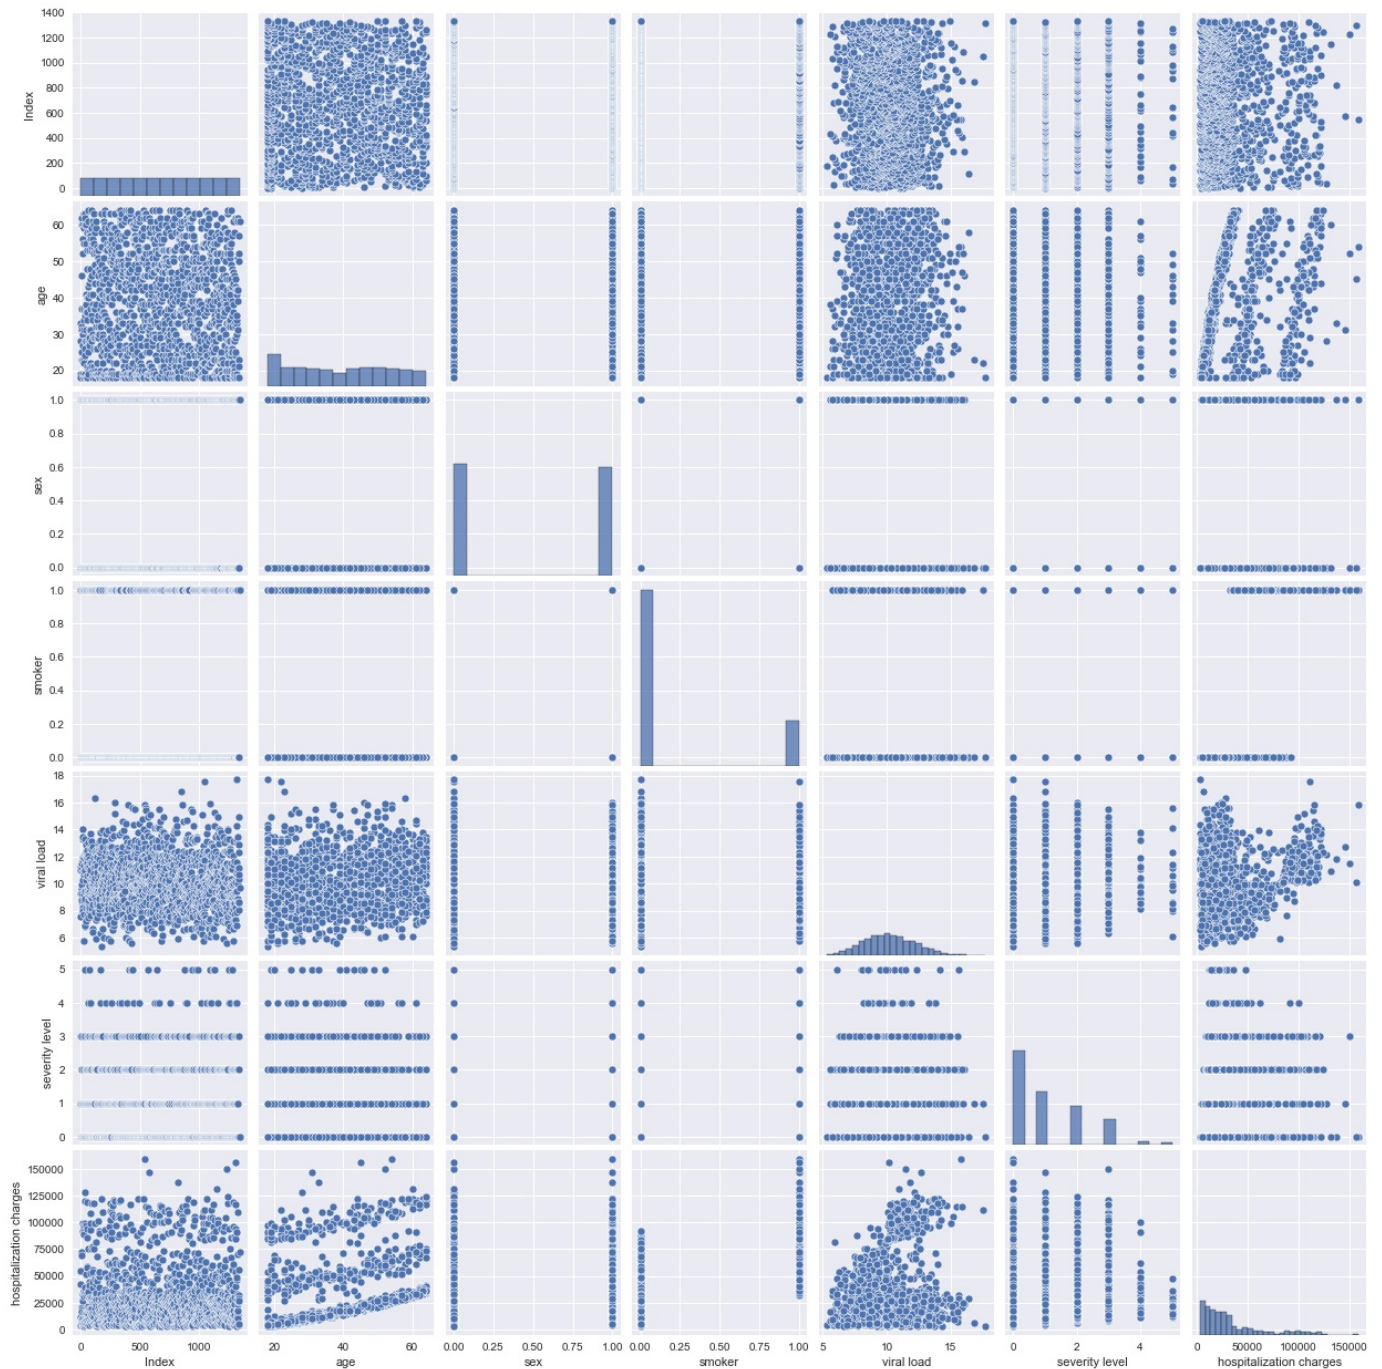
#Hospitalization charges and age
#Hospitalization charges and smoker
# Both show strong positive correlation
```

Out[39]: <AxesSubplot:>



```
In [40]: sns.pairplot(df)
```

```
Out[40]: <seaborn.axisgrid.PairGrid at 0x2c3475f9880>
```



```
In [41]: #Viral load by sex
df.groupby('sex')['viral load'].mean()
```

```
Out[41]: sex
0      10.314423
1      10.126073
Name: viral load, dtype: float64
```


```
In [42]: #hospitalization by smoking
df.groupby('smoker')['hospitalization charges'].mean()

#the difference for hospitalization charges for smoker and non smoker is significant
```

```
Out[42]: smoker
0      21085.675752
1      80125.572993
Name: hospitalization charges, dtype: float64
```

```
In [43]: #hospitalization charges by severity level
df.groupby('severity level')['hospitalization charges'].mean().plot()
```

```
Out[43]: <AxesSubplot:xlabel='severity level'>
```



```
In [44]: #viral load by smoking category
df.groupby('smoker')['viral load'].mean()


#viral load is independent of whether patient is smoker or not as per this
```

```
Out[44]: smoker
0      10.217378
1      10.236204
Name: viral load, dtype: float64
```

```
In [45]: df.groupby('age')['hospitalization charges'].mean().plot()

#strongly correlated
```

```
Out[45]: <AxesSubplot:xlabel='age'>
```




```
In [46]: #grouping age into bins
bins = [0,18,28,38,48,58,68]
labels = ['0-18','18-28','28-38','38-48','48-58','58-68']
df['age bins'] = pd.cut(df['age'], bins, labels)
```

```
In [47]: df.groupby('age bins')['hospitalization charges'].mean().plot()

#as age increases we see hospitalization charges increase linearly
```


Out[47]: <AxesSubplot:xlabel='age bins'>



In [48]:

```
plt.figure(figsize = (12,8))
sns.barplot(x='age bins', y = 'hospitalization charges', data = df)
```

Out[48]: <AxesSubplot:xlabel='age bins', ylabel='hospitalization charges'>





In [49]:

```
plt.figure(figsize = (10,8))
sns.lineplot(x='age', y = 'hospitalization charges', hue = 'sex', data = df)
```




#not muhc of a difference between males and females in hospitalization charges vs age.

Out[49]: <AxesSubplot:xlabel='age', ylabel='hospitalization charges'>





```
In [53]: plt.figure(figsize = (15,10))
cat_cols = ['sex', 'smoker', 'severity level']
for i,j in enumerate(cat_cols):
    #row 1, col 3, index begins i is 1 based hence i+1
    plt.subplot(1,3,i+1)
    plt.subplots_adjust(wspace = 0.5, hspace = 2)
    sns.boxplot(x=j, y='hospitalization charges', data=df)
```




```
In [51]: #Except for smokers the hospitalization charges do not change for across sex, severity level and region

plt.figure(figsize = (15,10))
sns.boxplot(x='region', y='hospitalization charges', data=df)
```

```
Out[51]: <AxesSubplot:xlabel='region', ylabel='hospitalization charges'>
```





```
In [52]: df.groupby(['sex','smoker','age bins'])['hospitalization charges'].mean().unstack()
```

	age bins	(0, 18]	(18, 28]	(28, 38]	(38, 48]	(48, 58]	(58, 68]
0	0	6739.285714	9617.187500	16996.948454	21754.786408	29646.357798	36771.500000
	1	61948.750000	76063.710526	74800.942857	83855.457143	91621.370370	107589.625000
1	0	9291.862069	13081.352000	14050.326923	22979.908257	31545.153226	39288.660714
	1	67156.000000	62991.233333	77482.782609	76219.586207	90458.214286	93530.866667

```
In [ ]:
```

```
In [55]: #Missing value treatment


#There are no missing values as we have seen earlier

#Outlier Treatment

# Hospitalization charges, Age and Viral Load are the variables that we need to check for outlier treatment


sns.boxplot(y=df['age'])
```


```
Out[55]: <AxesSubplot:ylabel='age'>
```



```
In [56]: sns.boxplot(y=df['hospitalization charges'])
```

```
Out[56]: <AxesSubplot:ylabel='hospitalization charges'>
```





In [57]: `sns.boxplot(y=df['viral load'])`

Out[57]: <AxesSubplot:ylabel='viral load'>



In [58]: *#Age does not have outliers as is clear from the boxplot and also the statistical summary drawn earlier*
#we treat viral load and hospitalization charges for outlier treatment


```
o_cols = ['viral load', 'hospitalization charges']


for i in o_cols:
    Q1ss = df[i].quantile(0.25)
    Q3ss = df[i].quantile(0.75)
    IQRss = Q3ss - Q1ss
    df = df[ ~ ((df[i] < (Q1ss - 1.5 * IQRss)) | (df[i] > (Q3ss + 1.5 * IQRss))) ]
```

In [59]: `df.shape`

Out[59]: (1191, 9)




In [62]: `for i,j in enumerate(o_cols):`
#row 1, col 3, index begins i is 1 based hence i+1
`plt.subplot(1,2,i+1)`
`plt.subplots_adjust(wspace = 0.5, hspace = 2)`
`sns.boxplot(y =df[j], data=df)`






In [59]: *#There are still outliers present but their range and number has reduced significantly. Applying outlier treatment again would lead to loss of data i.e. smaller sample size as more rows would be deleted*

```
In [65]: plt.figure(figsize = (15,10))
for i,j in enumerate(cat_cols):
    #row 1, col 3, index begins i is 1 based hence i+1
    plt.subplot(1,3,i+1)
    plt.subplots_adjust(wspace = 0.5, hspace = 2)
    sns.boxplot(x=j, y='hospitalization charges', data=df)
```



```
In [66]: plt.figure(figsize = (15,10))
sns.boxplot(x='region', y='hospitalization charges', data=df)
```

Out[66]: <AxesSubplot:xlabel='region', ylabel='hospitalization charges'>





```
In [67]: #Hypothesis Testing
```

```
In [68]: #Prove (or disprove) that the hospitalization charges of people who do smoking are greater than those who don't?
#We can do it using right tailed t test

smog = df.groupby('smoker')
```

```
In [69]: smog['hospitalization charges'].describe()
```

```
Out[69]:
```

	count	mean	std	min	25%	50%	75%	max
smoker								
0	1055.0	20907.971564	14563.067125	2805.0	9962.5	18313.0	28387.5	83680.0
1	136.0	54578.154412	13360.849267	32074.0	44663.5	51899.5	61421.5	85758.0

```
In [70]: smoker = df[df['smoker'] == 1]['hospitalization charges'].sample(136)
non_smoker = df[df['smoker'] == 0]['hospitalization charges'].sample(136)
```

```
In [71]: smoker.reset_index(drop=True)
non_smoker.reset_index(drop=True)
```

```
Out[71]:
```

0	4262
1	15281
2	12210
3	11072
4	32660
...	...
131	17905
132	5756
133	31608
134	26761
135	24290

Name: hospitalization charges, Length: 136, dtype: int64

```
In [67]: # 1) Prove (or disprove) that the hospitalization charges of people who do smoking are greater than those who don't
```

```
In [72]: # We use Right Tailed t test

# Null hypothesis : The mean of hospitalization charges of smokers are non smokers is equal
# Alternate hypothesis : The mean of hospitalization charges of smokers is greater than non smokers
#Checkin the assumptions

from scipy import stats
import statsmodels.api as sm
from statsmodels.graphics.gofplots import qqplot
from scipy.stats import norm, uniform
from scipy.stats import levene

stats.shapiro(smoker)
```


```
Out[72]: ShapiroResult(statistic=0.9291411638259888, pvalue=2.4488156213919865e-06)
```

```
In [69]: stats.shapiro(non_smoker)
```


```
Out[69]: ShapiroResult(statistic=0.8702459931373596, pvalue=1.4824373950617087e-09)
```

```
In [70]: #For both smoker and non smoker the p value are much smaller than 0.05 hence we know that they are NOT normally distributed. We further check this with q-q plot
```

```
In [71]: qqplot(smoker,norm,fit=True,line="45")  
plt.show()
```



```
In [72]: qqplot(non_smoker,norm,fit=True,line="45")  
plt.show()
```



```
In [73]: #Standard deviations  
np.var(smoker)
```

```
Out[73]: 177199702.73350984
```

```
In [74]: np.var(non_smoker)
```

```
Out[74]: 280250246.6286765
```

```
In [73]: #Std deviations are not equal. We can use levene's test to confirm . Levene's test gives a YES because it takes a certain factor by which the standard deviations can differ to be considered equal  
alpha = 0.05  
w_stats, p_value = levene(smoker, non_smoker, center = 'median')
```

```
In [74]: if p_value > alpha :
        print("Variances of smokers and non_smokers are equal")
        else:
        print("Variances of smokers and non_smokers are NOT equal")
```

Variances of smokers and non_smokers are equal

```
In [76]: stats.shapiro(np.log(smoker))
```

```
Out[76]: ShapiroResult(statistic=0.9721938967704773, pvalue=0.00695952819660306)
```

```
In [77]: stats.shapiro(np.log(non_smoker))
```

```
Out[77]: ShapiroResult(statistic=0.9450490474700928, pvalue=3.225280670449138e-05)
```

```
In [78]: #Even the log of smoker and non_smoker aren't normally distributed
```

```
In [83]: #For smoker and non_smoker
#1) The distributiùons are not normal for both of them
#2) But their variances are equal
```

#Proceeding with the T test

```
def t_test(x,y,alternative):
    _, double_p = stats.ttest_ind(x,y)

    if alternative == 'two-sided':
        pval = double_p
    elif alternative == 'greater':
        if np.mean(x) > np.mean(y):
            pval = double_p/2.
        else:
            pval = 1.0 - double_p/2.
    elif alternative == 'less':
        if np.mean(x) < np.mean(y):
            pval = double_p/2.
        else:
            pval = 1.0 - double_p/2.
    return pval
```

```
In [86]: p_val = t_test(smoker, non_smoker, 'greater')
        print(p_val)
```

2.8995676943672616e-63

```
In [87]: stats.ttest_ind(smoker, non_smoker, alternative='greater')
```

```
Out[87]: Ttest_indResult(statistic=22.23571094828444, pvalue=2.8995676943672616e-63)
```

```
In [88]: #As we can see on calcaulating from both ways p value is smaller than alpha = 0.05 . REJECT NULL HYPOTHESIS

if p_val < 0.05:
    print("We reject the null hypothesis ")
    print('"'The means hospitalization charges for smokers and non smokers are not equal. The mean hospitalization
charges for smokers and greater than non smokers'"')
else:
    print("The means hospitalization charges for smokers and non smokers are equal")
```

We reject the null hypothesis

The means hospitalization charges for smokers and non smokers are not equal. The mean hospitalization charges for smokers and greater than non smokers

```
In [ ]:
```

```
In [89]: #2) Prove (or disprove) with statistical evidence that the viral load of females is different from that of males

#We use two sided T test

# Null hypothesis : The mean (of) viral load of males and females is the same.

# Alternate hypothesis : The mean (of) viral load of males and females are unequal i.e. not the same.

# Checkin the assumptions
```

```
In [90]: df.groupby('sex')['viral load'].describe()
```

```
Out[90]:
```

	count	mean	std	min	25%	50%	75%	max
sex								
0	581.0	10.030947	1.94495	5.32	8.600	9.940	11.2900	15.51
1	610.0	9.966541	1.96940	5.60	8.585	9.855	11.1725	15.58

```
In [105]: # male viral load
mvl = df[ df['sex'] == 0]['viral load'].sample(581)

# female viral load
fvl = df[ df['sex'] == 1]['viral load'].sample(581)

#viral load irrespective of gender
tvL = df['viral load']
```

```
In [106]: mvl.reset_index(drop=True)
fvl.reset_index(drop=True)
```

```
Out[106]:
```

0	9.47
1	12.28
2	9.97
3	14.13
4	7.22
...	
576	9.70
577	10.54
578	11.97
579	5.76
580	9.69

Name: viral load, Length: 581, dtype: float64

```
In [107]: stats.shapiro(mvl)
```

```
Out[107]: ShapiroResult(statistic=0.9929239749908447, pvalue=0.007620655465871096)
```

```
In [108]: stats.shapiro(fvl)
```


```
Out[108]: ShapiroResult(statistic=0.9918439388275146, pvalue=0.002799693727865815)
```


```
In [109]: stats.shapiro(tvL)
```

```
Out[109]: ShapiroResult(statistic=0.9924992918968201, pvalue=9.647324986872263e-06)
```


```
In [110]: #For both male viral load and female viral load the p value are much smaller than alpha 0.05 hence we know that t
#are NOT normally distributed. We further check this with q-q plot

qqplot(mvl,norm,fit=True,line="45")
plt.show()
```






```
In [111]: qqplot(fvl,norm,fit=True,line="45")
plt.show()
```



```
In [112]: qqplot(tvl,norm,fit=True,line="45")
plt.show()
```



```
In [98]: #From q-q- plots we can see that male as well as female viral loads are not normally distributed. Viral load irre
# gender is also not normally distributed
```

```
In [113]: stats.shapiro(np.log(mvl))
```

```
Out[113]: ShapiroResult(statistic=0.9944441914558411, pvalue=0.03279054909944534)
```

```
In [114]: stats.shapiro(np.log(fvl))
```

Out[114... ShapiroResult(statistic=0.9927094578742981, pvalue=0.0062293680384755135)

```
In [ ]: #Even their logs are not normally distributed hence this isn't a log normal distribution
```

```
In [189... # We use levene's test to check for equal variances  
w_stats, p_value = levene(mvl, fvl, center="median")
```

```
In [190... if p_value > alpha :  
    print("Variances of viral loads of males and females are equal")  
else:  
    print("Variances of viral loads of males and females are NOT equal")
```

Variances of viral loads of males and females are equal

```
In [105... #Variances are equal.. lets confirm by calculating the variances individually as well  
  
np.var(mvl)
```

Out[105... 3.776321134846738

```
In [106... np.var(fvl)
```

Out[106... 3.9323172007429767

```
In [109... #Two sided T test on male viral load and female viral load  
  
t_test(mvl, fvl, 'two-sided')
```

Out[109... 0.6826719822745506

```
In [110... stats.ttest_ind(mvl, fvl, alternative='two-sided')
```

Out[110... Ttest_indResult(statistic=0.4089225510074076, pvalue=0.6826719822745506)

```
In [ ]: #p value is 0.68 hence we fail to reject the Null hypothesis
```

```
In [116... #Is the proportion of smoking significantly different across different regions?  
  
#We use a Chi squared test  
  
#Null hypothesis (H0) : Smoker and regions are independent variables i.e. mutually independent. There is no relation  
  
#Alternate hypothesis (H1) : Both these variables are dependent  
  
contingency= pd.crosstab(df['region'], df['smoker'])  
contingency
```

Out[116...

smoker	0	1
region		
northeast	256	39
northwest	267	38
southeast	267	34
southwest	265	25

```
In [117... from scipy.stats import chi2_contingency  
  
stat, p, dof, expected = chi2_contingency(contingency, correction = False)
```



```
In [118... p
```

Out[118... 0.31791538258247426

```
In [119... if p <= alpha:
    print('Dependent (reject H0)')
else:
    print('Independent (Fail to reject H0)')

Independent (Fail to reject H0)
```

```
In [ ]:
```

```
In [140... #Is the mean viral load of women with 0 Severity level, 1 Severity level, and 2 Severity level the same?
# H0 (Null Hypothesis): The mean viral load of women with 0 Severity level, 1 Severity level, and 2 Severity level are the same
# H1 (Alternate Hypothesis) : The mean viral loads of women with severity level 0, 1 and 2 are not the same

df.groupby(['sex', 'severity level'])['viral load'].mean()
```

Out[140... sex severity level

0	0	10.034880
	1	10.042113
	2	10.026162
	3	10.023881
	4	10.239231
	5	9.598000
1	0	9.963209
	1	9.908844
	2	9.945000
	3	10.014366
	4	10.601000
	5	10.206250

Name: viral load, dtype: float64

```
In [120... #We can observe that the mean viral load of women with 0 Severity level, 1 Severity level, and 2 Severity level are the same
#Confirming the same with Anova
```

```
In [121... df[df['sex'] == 1].groupby('severity level')['viral load'].describe()
```

Out[121...

		count	mean	std	min	25%	50%	75%	max
severity level									
	0	268.0	9.963209	1.936862	5.76	8.6075	9.695	11.1300	14.92
	1	147.0	9.908844	1.918987	5.60	8.6000	9.670	11.1300	15.36
	2	106.0	9.945000	2.092305	5.73	8.3450	10.060	11.2025	15.57
	3	71.0	10.014366	1.950361	6.33	8.5950	10.030	11.1750	14.90
	4	10.0	10.601000	1.815063	8.53	9.6200	9.825	11.0350	13.82
	5	8.0	10.206250	2.975480	6.10	8.0500	10.080	11.6625	15.58

```
In [151... dfws = df[(df['sex'] == 1) & (df['severity level'] <=2)]
```

```
In [154... dfws.reset_index(drop=True)
```

Out[154...

	Index	age	sex	smoker	region	viral load	severity level	hospitalization charges	age bins
0	0	19	1	1	southwest	9.30	0	42212	(18, 28]
1	5	31	1	0	southeast	8.58	0	9392	(28, 38]
2	6	46	1	0	southeast	11.15	1	20601	(38, 48]
3	9	60	1	0	northwest	8.61	0	72308	(58, 68]
4	11	62	1	1	southeast	8.76	0	69522	(58, 68]

...
516	1331	23	1	0	southwest	11.13	0	26990	(18, 28]
517	1334	18	1	0	northeast	10.64	0	5515	(0, 18]
518	1335	18	1	0	southeast	12.28	0	4075	(0, 18]
519	1336	21	1	0	southwest	8.60	0	5020	(18, 28]
520	1337	61	1	1	northwest	9.69	0	72853	(58, 68]

521 rows × 9 columns

```
In [178... #Checking assumption of Anova that each group sample is drawn from a normally distributed population
#for females taking into account the severity level 0,1, and 2 the viral load isn't normally distributed
stats.shapiro(dfws['viral load'])
```

```
Out[178... ShapiroResult(statistic=0.9920216202735901, pvalue=0.006814346183091402)
```

```
In [180... #for females taking into account all severity levels the viral load isn't normally distributed
stats.shapiro(df[df['sex'] == 1]['viral load'])
```


```
Out[180... ShapiroResult(statistic=0.9914260506629944, pvalue=0.0013309171190485358)
```

```
In [183... #female viral load for severity level 0,1 and 2 is not log normally distributed
stats.shapiro(np.log(dfws['viral load']))
```


```
Out[183... ShapiroResult(statistic=0.9924635291099548, pvalue=0.009911485947668552)
```


```
In [181... #We further confirm this qq plot

qqplot(dfws['viral load'],norm,fit=True,line="45")
plt.show()
```



```
In [182... qqplot(df[df['sex'] == 1]['viral load'],norm,fit=True,line="45")
plt.show()
```





```
In [ ]: #As is clear from the qq plot
# 1) The viral load for females with severity level 0,1 and 2 is not normally distributed
# 2) The viral load for females with taking into account all severity levels isn't normally distributed
```

```
In [186... dws2 = []
for i in range(0,3):
    dws2i = df[(df["sex"] == 1) & (df['severity level'] == i)]["viral load"].sample(106)
    dws2.append(dws2i)
```

```
In [191... #Checking for assumptions that all groups have same variance

w_stats, p_value = levene(dws2[0], dws2[1], dws2[2], center = 'mean')
```

```
In [192... if p_value > alpha :
    print("Variances of viral load 0,1 and 2 are equal in females")
else:
    print("Variances of viral load 0,1 and 2 are NOT Equal in females")
```

Variances of viral load 0,1 and 2 are equal in females

```
In [194... from scipy.stats import f_oneway

fstat, p_value = f_oneway(dws2[0], dws2[1], dws2[2])
```

```
In [196... p_value
```

```
Out[196... 0.9561554023268642
```

```
In [197... if p_value > alpha :
    print("Means of viral load 0,1 and 2 are equal in females. We accept the null hypothesis")
else:
    print("Means of viral load 0,1 and 2 are NOT Equal in females and we fail to reject the null hypothesis")
```

Means of viral load 0,1 and 2 are equal in females. We accept the null hypothesis

```
In [122... df.groupby('region')['hospitalization charges'].mean()
```

```
Out[122... region
northeast    26851.705085
northwest    25908.986885
southeast    23854.235880
southwest    22334.206897
Name: hospitalization charges, dtype: float64
```

```
In [ ]:
```